

CIS 6930 Special Topics in Large Language Models

N-gram Language Models

Language Models (LM)

- LM should have the ability to assign each possible sequence of tokens $w_1, \dots, w_n \in V$ with a meaningful probability: $p(w_1, \dots, w_n)$

$p(w_1, \dots, w_n)$	Joint Probability	Chain Rule
$= p(w_1) * p(w_2 w_1) * p(w_3 w_1, w_2) * \dots * p(w_n w_1, w_2, \dots, w_{n-1})$		

E.g. $p(\text{cat sat on the mat}) = p(\text{cat}) * p(\text{sat} | \text{cat}) * p(\text{on} | \text{cat sat}) * p(\text{the} | \text{cat sat on}) * p(\text{mat} | \text{cat sat on the})$

Conditional Probability

Language Models (LM)

- LM should have the ability to assign each possible sequence of tokens $w_1, \dots, w_n \in V$ with a meaningful probability: $p(w_1, \dots, w_n)$
- LM is a machine learning model that assigns a probability for each possible next word $p(w_n | w_1, \dots, w_{n-1})$ (predicting upcoming words)
- This is how LLM work: (1) LLMs are **trained** to predict words (2) LLMs **generate** text by predicting the next word repeatedly.

N-gram Language Model

- N-gram Language Model – the simplest kind of language model
- What is n-gram: n-gram is a sequence of n words
 - 2-gram (bigram) = two-word sequence of words, e.g. the cat, cat sat...
 - 3-gram (trigram) = three-word sequence of words, e.g. the cat sat ...
- N-gram LM is a probabilistic model that can estimate the prob of a word given the **N-1** previous words, and thereby to assign prob for entire sequences

Outline

- Building N-gram Language Model
- Generation from N-gram Language Model
- Evaluating N-gram Language Model
- Smoothing of N-gram Language Model

Building a N-gram Language Model

How to estimate probability

- One way to estimate $p(w|h)$ prob of a word w given some history h is using **relative frequency counts**

$$P(\text{sat}|\text{the cat}) = \frac{\text{count}(\text{the cat sat})}{\text{count}(\text{the cat})}$$

$$P(\text{on}|\text{the cat sat}) = \frac{\text{count}(\text{the cat sat on})}{\text{count}(\text{the cat sat})}$$

MLE: Maximum Likelihood Estimation

Problems of relative frequency counts

- **Not accurate** because new sentences are invented all the time
Even the entire web is not big enough to give us good estimates for counts of every possible sentences
- **Computational heavy** for sequence with long length n
With a vocabulary of size V , # sequence of length n = V^n
Typical English vocabulary $V \sim 40k$ words
Even sentences of length ≤ 11 results in more than 4×10^{50} sequences

Markov Assumption

- Markov Assumption definition – the future state conditioned on the current state (or recent several states) is independent of past states
- Use only the recent past to predict the next word

1-st order

$$P(\text{mat}|\text{the cat sat on the}) \approx P(\text{mat}|\text{the})$$

2-nd order

$$P(\text{mat}|\text{the cat sat on the}) \approx P(\text{mat}|\text{on the})$$

N-gram Language Model

- Instead of computing the prob of word given its entire history, we only use the last **N-1** steps to approximate all the history

$$P(w_n | w_{1:n-1}) \approx P(w_n | w_{n-N+1:n-1})$$

$$p(w_1, \dots, w_n) = \prod_i p(w_i | w_{i-N+1}, \dots, w_{i-1})$$

N-gram Language Model

Unigram

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i)$$

e.g. P(the) P(cat) P(sat)

Bigram

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1})$$

e.g. P(the) P(cat | the) P(sat | cat)

$$p(w_i | w_{i-1}) = \frac{C(w_{i-1}, w_i)}{C(w_{i-1})}$$

Larger N , more accurate and better the LM is, but also higher costs

Practice – Estimating Conditional Probabilities

Consider the following corpus

< s > I like apples < /s >

< s > You like strawberries < /s >

< s > You like apples < /s >

Note: < s > and < /s > are starting and ending tokens

What's the bigram probability $P(\text{apples} \mid \text{like})$?

- (A) 1/3
- (B) 2/3
- (C) 1/2
- (D) 1

Practice – Estimating Conditional Probabilities

Consider the following corpus

< s > I like apples < /s >

< s > You like strawberries < /s >

< s > You like apples < /s >

Note: < s > and < /s > are starting and ending tokens

What's the bigram probability $P(\text{apples} \mid \text{like})$?

- (A) 1/3 (B) 2/3 (C) 1/2 (D) 1

$$P(\text{apples} \mid \text{like}) = \frac{\text{Count}(\text{"like apples"})}{\text{Count}(\text{"like"})} = \frac{2}{3}$$

Practice – Estimating Joint Probabilities

Consider the following corpus

< s > I like apples < /s >

< s > You like strawberries < /s >

< s > You like apples < /s >

Note: < s > and < /s > are starting and ending tokens

Using the bigram model, what's the probability of the sentence “< s > I like strawberries < /s >”? Ignore the probability of < s >.

- (A) 4/9
- (B) 1/3
- (C) 2/9
- (D) 1/9

Practice – Estimating Joint Probabilities

Consider the following corpus

< s > I like apples < /s >

< s > You like strawberries < /s >

< s > You like apples < /s >

Note: < s > and < /s > are starting and ending tokens

Using the bigram model, what's the probability of the sentence “< s > I like strawberries < /s >”? Ignore the probability of < s >.

- (A) 4/9 (B) 1/3 (C) 2/9 (D) 1/9

$$P(< s > \text{ I like strawberries } < /s >) = \frac{1}{3} \cdot 1 \cdot \frac{1}{3} \cdot 1$$



Scaling in Large N-gram Models – Log Probability

- Probability Numerical Underflow: Because probabilities are ≤ 1 , the more probs we multiply together, the smaller the product becomes
- Log Probability: adding in log space is equivalent to multiplying in linear space, and we get results that are not small

$$p_1 \times p_2 \times p_3 \times p_4 = \exp(\log p_1 + \log p_2 + \log p_3 + \log p_4)$$

Generation from N-gram Language Model

Sampling a word from a prob distribution



Generating sequences from bigram LM

$$\text{Bigram} \quad P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1})$$

- Generate the first word $w_1 \sim P(w)$
- Generate the second word $w_2 \sim P(w | w_1)$
- Generate the third word $w_3 \sim P(w | w_2)$
- ...

Generating sequences from bigram LM

Choose a random bigram ($< s >$, w)

according to its probability $p(w | < s >)$

Now choose a random bigram (w, x)
according to its probability $p(x | w)$

And so on until we choose $< /s >$

Then string the words together

$< s >$

I

I want

want to

to eat

eat Chinese

Chinese food

food

$< /s >$

I want to eat Chinese food

Generating sequences from trigram LM

$$\text{Trigram} \quad P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i \mid w_{i-2}, w_{i-1})$$

- Generate the first word $w_1 \sim P(w)$
- Generate the second word $w_2 \sim P(w \mid w_1)$
- Generate the third word $w_3 \sim P(w \mid w_1, w_2)$
- Generate the fourth word $w_4 \sim P(w \mid w_2, w_3)$
- ...

Evaluating N-gram Language Model


Two types of evaluation for LM

- Extrinsic Evaluation

Train LM -> apply to downstream task -> evaluate the task

Higher task performance -> better model

Problem: expensive, time consuming, indirect feedback



Two types of evaluation for LM

- Intrinsic Evaluation

Measures the quality of a LM independent of any application

How well a LM estimates the probability of sequences

How to evaluate: training set, testing set, development set, metric

Intrinsic Evaluation – split datasets

- **Training set:** is the data we use to learn the parameter of model

For the N-gram LM, training data is the data from which we get the n-gram counts to estimate n-gram probability

- **Testing set:** is a set of unseen data that is not overlapping with training set, that we use to evaluate the model

If a LM assigns a higher probability to the test set, that means it more accurately predicts the test set, then it is a better model.

- **Development set:** used to tune hyper-parameter

Intrinsic Evaluation – metric

- Recall: if a LM assigned a higher probability to the test set, it is better
- Question: Can we use probability as the evaluation metric?

We do not use the raw probability as the metric, because the prob of a test set depends on the number of words/tokens – the prob gets smaller if the test text is longer.
- Metric: **Perplexity** (a metric normalized by sequence length, per-word)

Intrinsic Evaluation – perplexity

- Perplexity has an inverse relation with probability
- Better LM = higher estimated probability on test set = lower perplexity

$$\begin{aligned}\text{perplexity}(W) &= P(w_1 w_2 \dots w_N)^{-\frac{1}{N}} \\ &= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}\end{aligned}$$

Chain Rule

$$\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$

Bigram LM

$$\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$


Intrinsic Evaluation – perplexity

- Example: unigram, bigram, trigram models on 38M words from the *Wall Street Journal* news paper, their perplexity on test corpus (1.5M words)

	Unigram	Bigram	Trigram
Perplexity	962	170	109

Trigram model gives us more information about word sequence, has a better idea of what words might come next, thus assigns a higher prob

Intrinsic Evaluation – perplexity



GPT-3 175B: perplexity = 20.5

Smoothing of N-gram Language Model

Problem of N-gram Language Model

- Not all n-grams in the test set will be observed in training data

Example: Training set is Google News, Testing set is Shakespeare

$$P(\text{affray} \mid \text{voice doth us}) = 0$$

- Problem of a particular n-gram never occurs in the training but appears in the test set:
 - (1) Underestimate the probability of words sequences that might occur
 - (2) The prob of the whole test set will also be 0, perplexity no sense

Smoothing

Smoothing is to make sure the probabilities for all the possible n-gram are non-zero in our model

- Laplace Smoothing: Adding a small amount to all probabilities
- Interpolation: Combining different granularities of n-gram LM
- Backoff: Backing off to lower n-gram

Smoothing – Laplace Smoothing

Adding a small amount to all probabilities

Max Likelihood Estimate for Bigram

$$P(w_i | w_{i-1}) = \frac{C(w_{i-1}, w_i)}{C(w_{i-1})}$$

After Laplace Smoothing

$$P(w_i | w_{i-1}) = \frac{C(w_{i-1}, w_i) + \alpha}{C(w_{i-1}) + \alpha |V|}$$

Smoothing – Linear Interpolation

Combining different granularities of n-gram LM to estimate probability

$$\hat{P}(w_i \mid w_{i-2}, w_{i-1}) = \lambda_1 P(w_i \mid w_{i-2}, w_{i-1}) \quad \text{Trigram}$$
$$+ \lambda_2 P(w_i \mid w_{i-1}) \quad \text{Bigram}$$
$$+ \lambda_3 P(w_i) \quad \text{Unigram}$$

$$\sum_i \lambda_i = 1$$

Process: (1) estimate n-gram prob on training set (2) tune hyper-parameter (lambdas) to maximize prob on dev set (3) evaluate on test set

Smoothing – Backoff

If the n-gram we need has zero counts, we approximate it by backing off to the (n-1) gram, until we reach a history that has non-zero counts.

Suppose you want:

$$P(\text{pancakes} \mid \text{delicious soufflé})$$

If the trigram probability is 0, use the bigram

$$P(\text{pancakes} \mid \text{soufflé})$$

If the bigram probability is 0, use the unigram

$$P(\text{pancakes})$$



Thank you!

UF | Herbert Wertheim
College of Engineering
UNIVERSITY *of* FLORIDA

LEADING THE CHARGE, CHARGING AHEAD

References

Speech and Language Processing (3rd ed. draft)

Dan Jurafsky and James H. Martin

Chapter 3