

# CIS 6930 Special Topics in Large Language Models

## Feedforward Neural Language Models

# Neural Networks in NLP

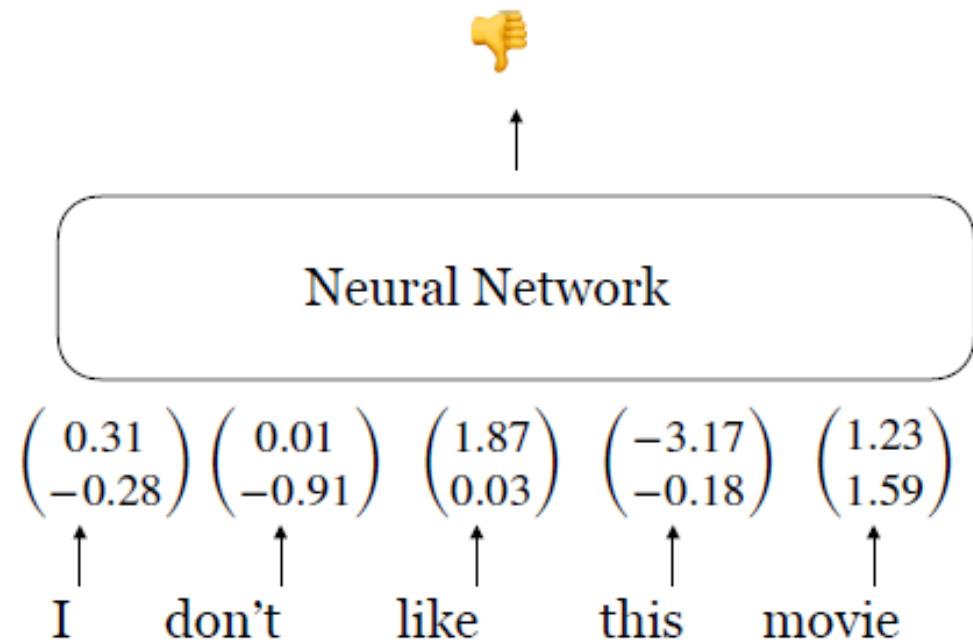
## From Word Embedding to Neural Networks

$$v_{\text{cat}} = \begin{pmatrix} -0.224 \\ 0.130 \\ -0.290 \\ 0.276 \end{pmatrix}$$

$$v_{\text{dog}} = \begin{pmatrix} -0.124 \\ 0.430 \\ -0.200 \\ 0.329 \end{pmatrix}$$

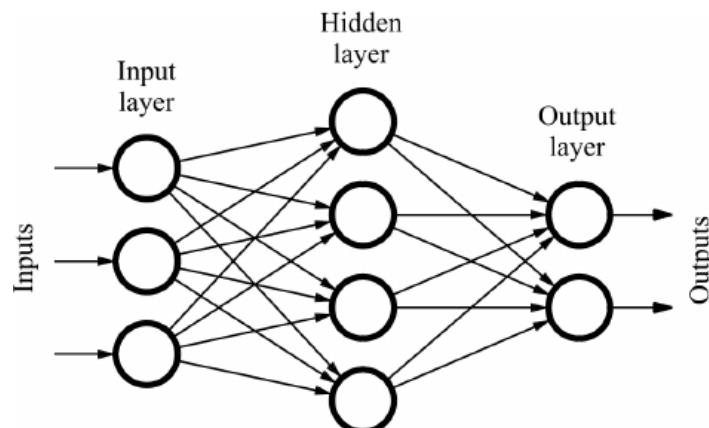
$$v_{\text{the}} = \begin{pmatrix} 0.234 \\ 0.266 \\ 0.239 \\ -0.199 \end{pmatrix}$$

$$v_{\text{language}} = \begin{pmatrix} 0.290 \\ -0.441 \\ 0.762 \\ 0.982 \end{pmatrix}$$

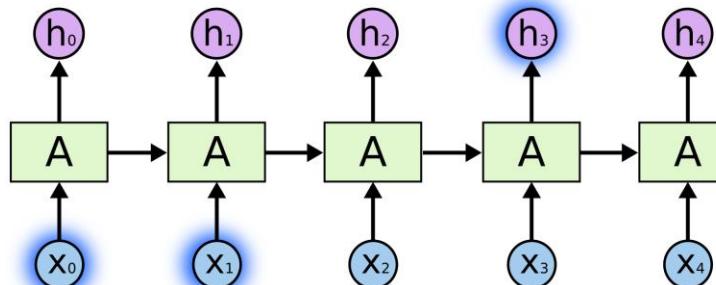


# Neural Networks in NLP

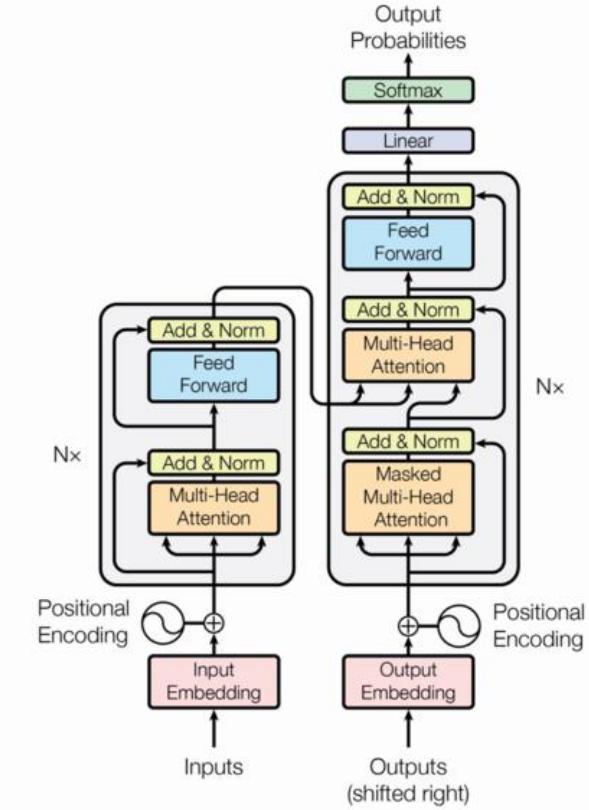
## Several Types of Neural Networks in NLP



Feed-forward NNs



Recurrent NNs



Transformers

# Outline

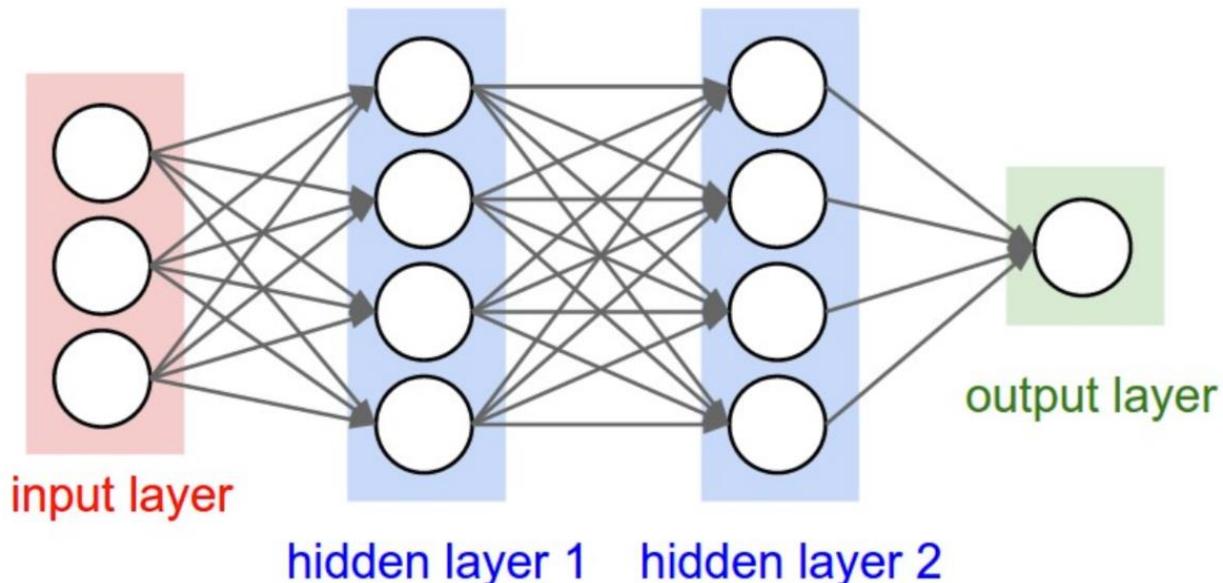
---

- Feedforward Neural Networks
- Model 1: Neural “bag-of-words” Models
- Model 2: Feedforward Neural Language Models

# **Feedforward Neural Networks**

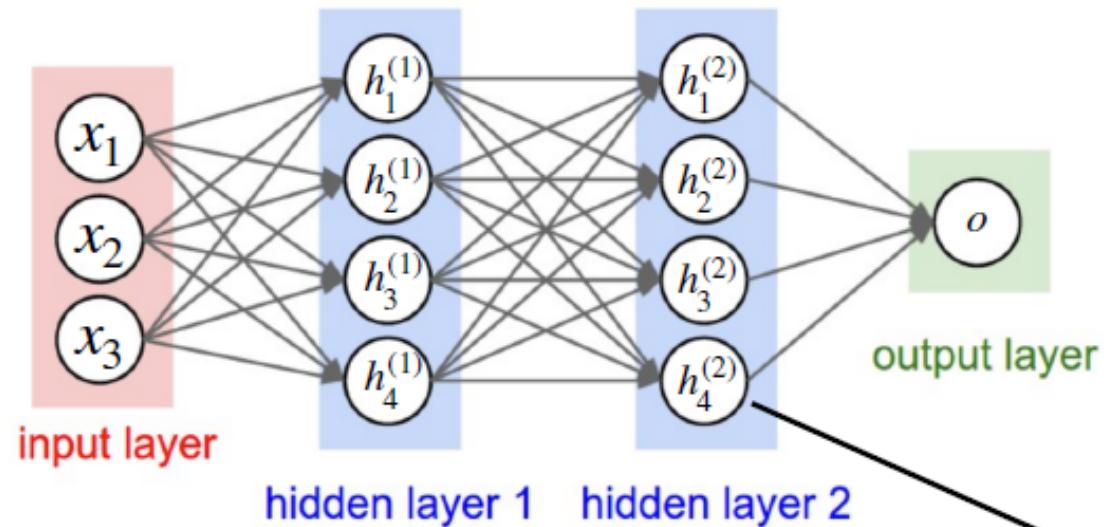
# Feed-forward Neural Network

- The units are connected with no cycles
- The outputs from units in each layer are passed to units in the next layer
- No outputs are passed back to lower layers



**Fully-connected (FC) layers**  
All the units from one layer are  
fully connected to every units  
of the next layer

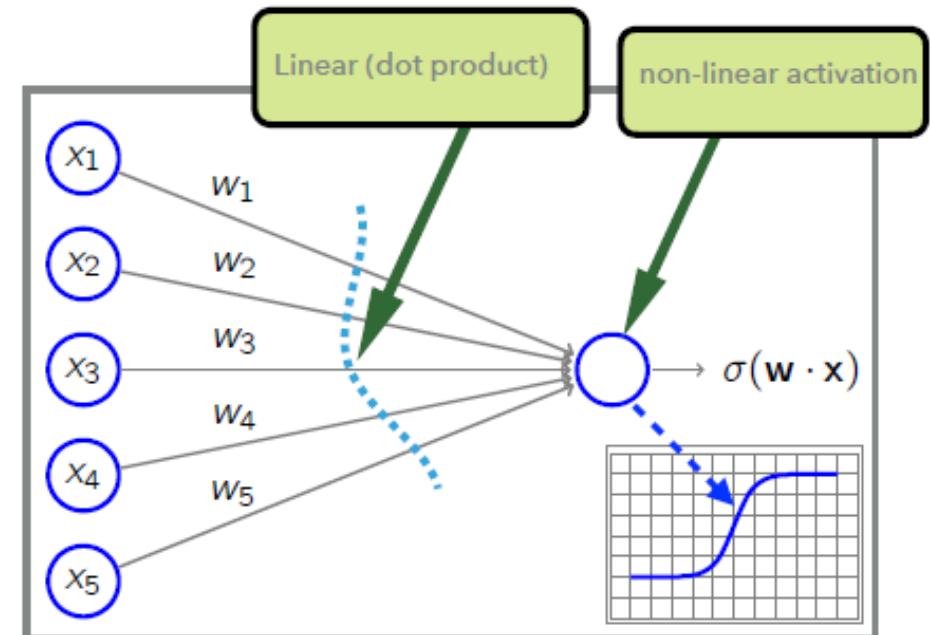
# Feed-forward Neural Network



$$h_1^{(1)} = f(w_{1,1}^{(1)}x_1 + w_{1,2}^{(1)}x_2 + w_{1,3}^{(1)}x_3)$$

$$h_3^{(2)} = f(w_{3,1}^{(2)}h_1^{(1)} + w_{3,2}^{(2)}h_2^{(1)} + w_{3,3}^{(2)}h_3^{(1)} + w_{3,4}^{(2)}h_4^{(1)})$$

non-linearity  $f$ :  $\sigma$ , tanh or ReLU.

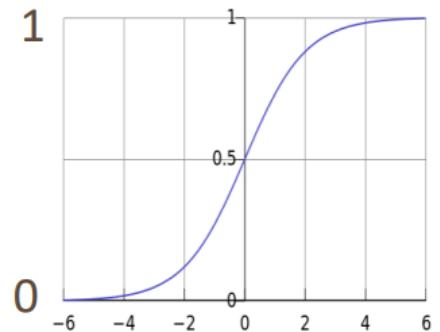


# Activation Functions

- For deep neural networks, the first thing to try is ReLU: it trains quickly and performs well due to good gradient backflow
- Sigmoid and tanh functions are still used (e.g. sigmoid to get probability)

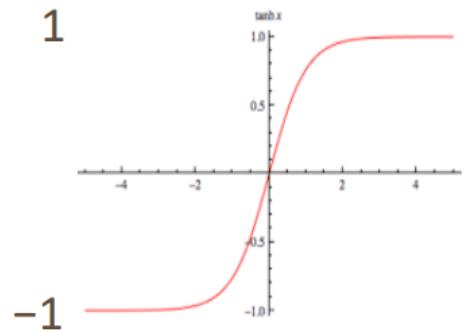
logistic (“sigmoid”)

$$f(z) = \frac{1}{1 + \exp(-z)}$$



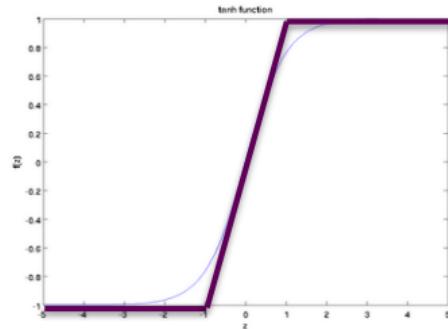
tanh

$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



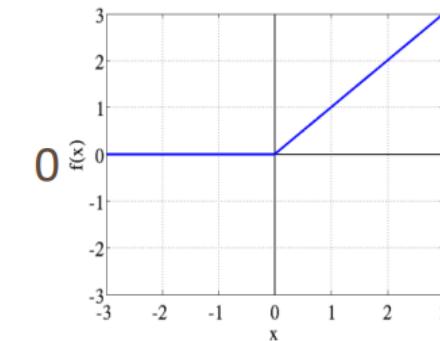
hard tanh

$$\text{HardTanh}(x) = \begin{cases} -1 & \text{if } x < -1 \\ x & \text{if } -1 \leq x \leq 1 \\ 1 & \text{if } x > 1 \end{cases}$$

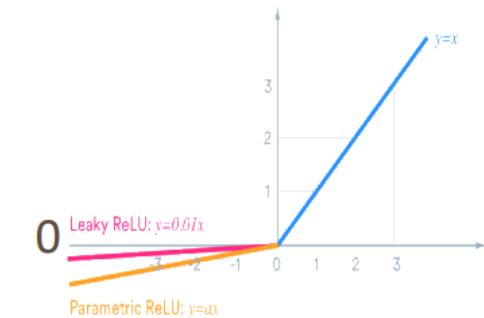


(Rectified Linear Unit)  
ReLU

$$\text{ReLU}(z) = \max(z, 0)$$

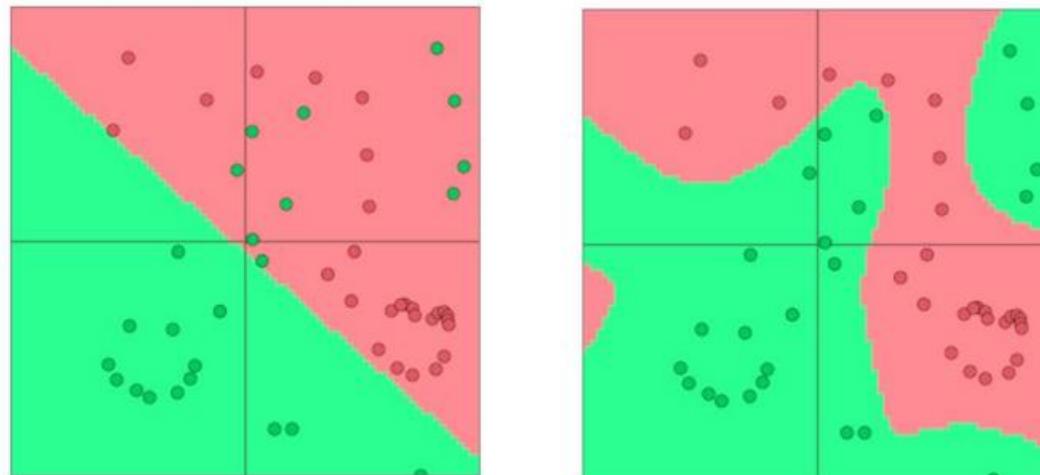


Leaky ReLU /  
Parametric ReLU

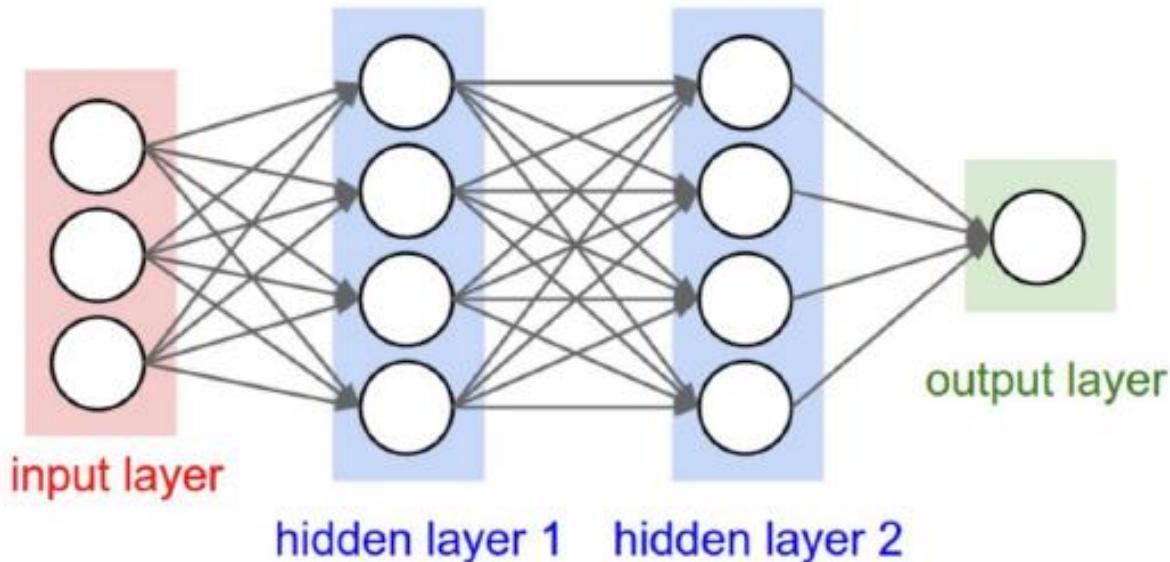


# Why Activation Functions

- The purpose is to add **Non-Linearities** into Neural Networks
- Neural networks do function approximation, without non-linearities, neural networks cannot do anything more than a linear transformation
- But with more layers that include non-linearities, neural networks can approximate any **complex** function!



# Matrix Notations



\*:  $f$  is applied element-wise

$$f([z_1, z_2, z_3]) = [f(z_1), f(z_2), f(z_3)]$$

$C$ : number of classes

$d$ : input dimension,  $d_1, d_2$ : hidden dimensions

- Input layer:  $\mathbf{x} \in \mathbb{R}^d$

- Hidden layer 1:

$$\mathbf{h}_1 = f(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) \in \mathbb{R}^{d_1}$$
$$\mathbf{W}^{(1)} \in \mathbb{R}^{d_1 \times d}, \mathbf{b}^{(1)} \in \mathbb{R}^{d_1}$$

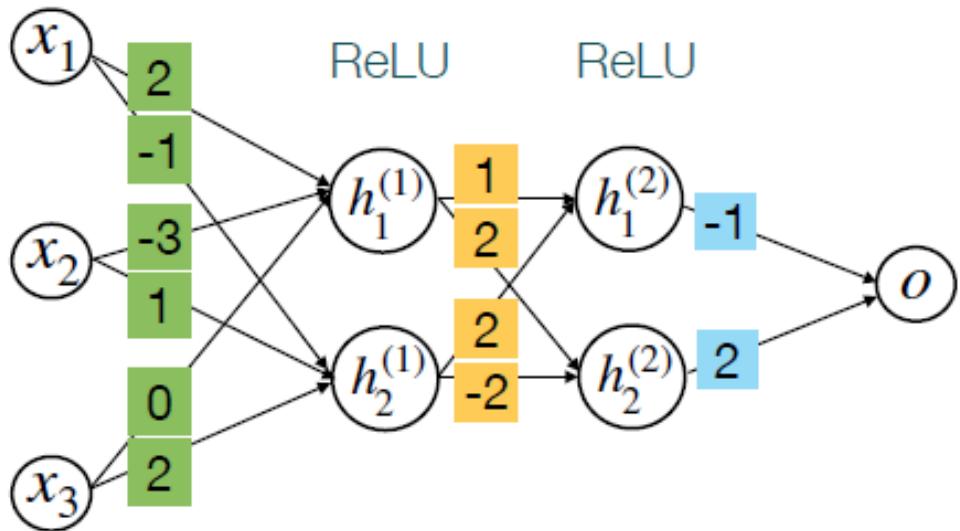
- Hidden layer 2:

$$\mathbf{h}_2 = f(\mathbf{W}^{(2)}\mathbf{h}_1 + \mathbf{b}^{(2)}) \in \mathbb{R}^{d_2}$$
$$\mathbf{W}^{(2)} \in \mathbb{R}^{d_2 \times d_1}, \mathbf{b}^{(2)} \in \mathbb{R}^{d_2}$$

- Output layer:

$$\mathbf{y} = \mathbf{W}^{(o)}\mathbf{h}_2, \mathbf{W}^{(o)} \in \mathbb{R}^{C \times d_2}$$

# Practice – Feedforward Neural Network



(Bias terms omitted in  
the next few slides)

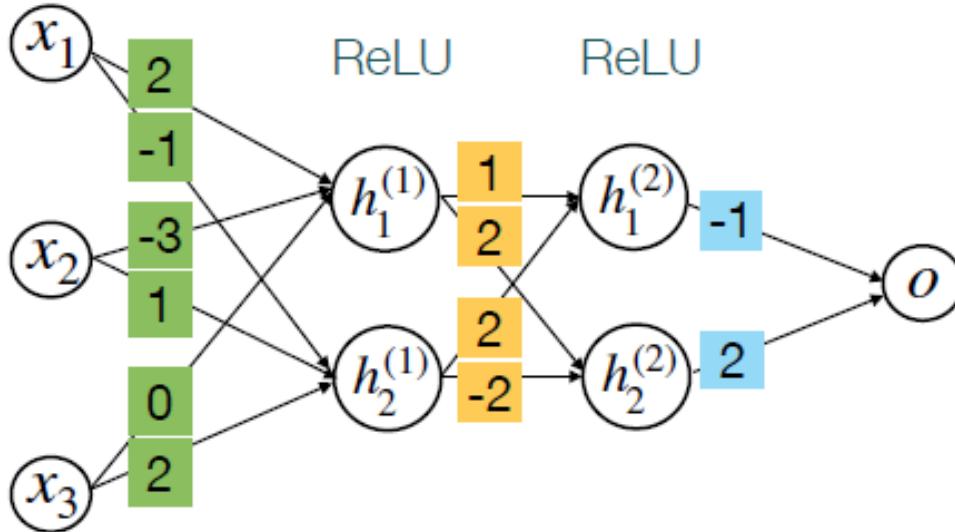
For  $x_1 = x_2 = x_3 = 1$ , what is the value of  $h_1^{(1)}$ ?

- (a) 0
- (b) -1
- (c) 1
- (d) 2

Correct: (a), because of the RELU:

$$\max(2 \times 1 + (-3) \times 1 + 0 \times 1, 0) = \max(-1, 0) = 0$$

# Practice – Feedforward Neural Network



(Bias terms omitted in  
the next few slides)

$$\mathbf{h}^{(1)} = \text{ReLU}(\mathbf{W}^{(1)}\mathbf{x}) \quad \text{What is the matrix } \mathbf{W}^{(1)}?$$

(a)  $\begin{bmatrix} 2 & -1 \\ -3 & 1 \\ 0 & 2 \end{bmatrix}$

(b)  $\begin{bmatrix} 2 & -3 & 0 \\ -1 & 1 & 2 \end{bmatrix}$

(c)  $\begin{bmatrix} 1 & 2 \\ 2 & -2 \end{bmatrix}$

Correct: (b).  $\mathbf{W}^{(1)}$  is a  $2 \times 3$  matrix.

# Training Feedforward Neural Networks

- Softmax: applying the softmax function on output logits to get predicted probability

$$\mathbf{y} = \mathbf{W}^{(o)} \mathbf{h}_2, \mathbf{W}^{(o)} \in \mathbb{R}^{C \times d_2}$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{y}) \quad \text{softmax}(\mathbf{y})_k = \frac{\exp(y_k)}{\sum_{j=1}^C \exp(y_j)} \quad \mathbf{y} = [y_1, y_2, \dots, y_C]$$

- Training loss: maximize the probability of the correct class  $y$  or equivalently we can minimize the negative log probability of that class

$$\min_{\mathbf{W}^{(1)}, \mathbf{W}^{(2)}, \mathbf{W}^{(o)}} - \sum_{(\mathbf{x}, y) \in D} \log \hat{\mathbf{y}}_y$$

# Training Feedforward Neural Networks

- Optimizer: stochastic gradient descent (SGD) to train feedforward NN

$$\theta^{new} = \theta^{old} - \alpha \nabla_{\theta} J(\theta)$$



stepping size or learning rate

Model parameters include  $W$  weight matrix and  $b$  bias terms between layers

Objective: learn model parameters to minimize loss

- Neural Networks are difficult to optimize. SGD can only converge to local minimum. **Initialization** and **Optimizer** matter a lot!

# Training Feedforward NN – Back-propagation

## Forward propagation:

from input to output layer

**Given:**  $x_1, x_2, x_3$   
and the class  
label  $y$   
(a single training  
example)

### Forward step 1:

Compute  $h_1^{(1)}, h_2^{(1)}$

### Forward step 2:

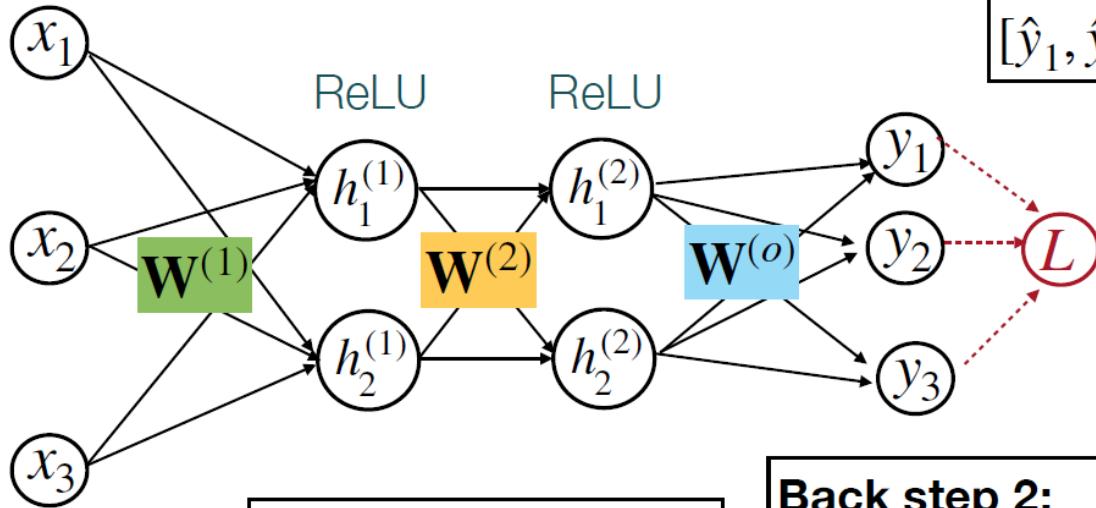
Compute  $h_1^{(2)}, h_2^{(2)}$

### Forward step 3:

Compute  $y_1, y_2, y_3$  and  
 $[\hat{y}_1, \hat{y}_2, \hat{y}_3] = \text{softmax}[y_1, y_2, y_3]$

### Forward step 4:

Compute loss  
 $L = -\log \hat{y}_y$



## Goal:

$$\frac{\partial L}{\partial W^{(1)}},$$

$$\frac{\partial L}{\partial W^{(2)}},$$

$$\frac{\partial L}{\partial W^{(o)}}$$

### Back step 4:

Compute  
 $\frac{\partial L}{\partial W^{(1)}}$

### Back step 3:

Compute  
 $\frac{\partial L}{\partial h_1^{(1)}}, \frac{\partial L}{\partial h_2^{(1)}}, \frac{\partial L}{\partial W^{(2)}}$

### Back step 2:

Compute  
 $\frac{\partial L}{\partial h_1^{(2)}}, \frac{\partial L}{\partial h_2^{(2)}}, \frac{\partial L}{\partial W^{(o)}}$

## Back propagation:

from output to input layer

# Training Feedforward NN – Back-propagation

```
1 import torch.nn as nn
2 import torch.nn.functional as F
3
4 class Net(nn.Module):
5     def __init__(self):
6         super().__init__()
7         self.fc1 = nn.Linear(784, 128)
8         self.fc2 = nn.Linear(128, 64)
9         self.fc3 = nn.Linear(64, 10)
10
11    def forward(self, x):
12        x = F.relu(self.fc1(x))
13        x = F.relu(self.fc2(x))
14        x = self.fc3(x)
15        return x
16
```

```
1 import torch.optim as optim
2
3 net = Net()
4 criterion = nn.CrossEntropyLoss()
5 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

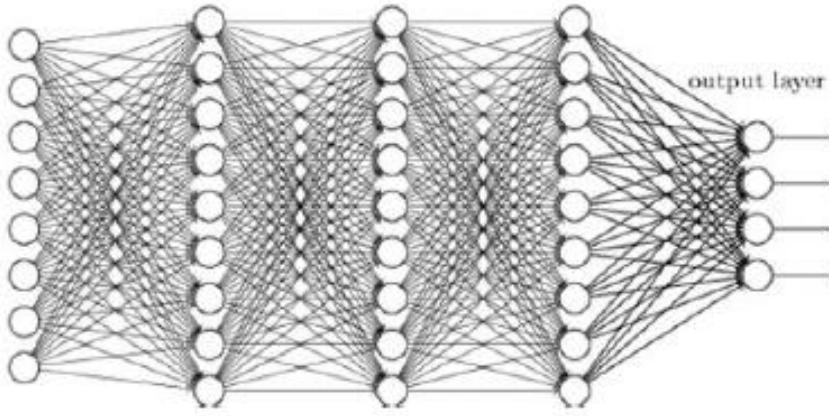
```
1 outputs = net(inputs)
2 loss = criterion(outputs, labels)
3 loss.backward()
4 optimizer.step()
```



*PyTorch did back-propagation for you in this one line of code!*

# **Neural “bag-of-words” Models for Text Classification**

# Comparison: image vs text inputs



label = “dog”

a sometimes tedious film

i had to look away - this was god awful .

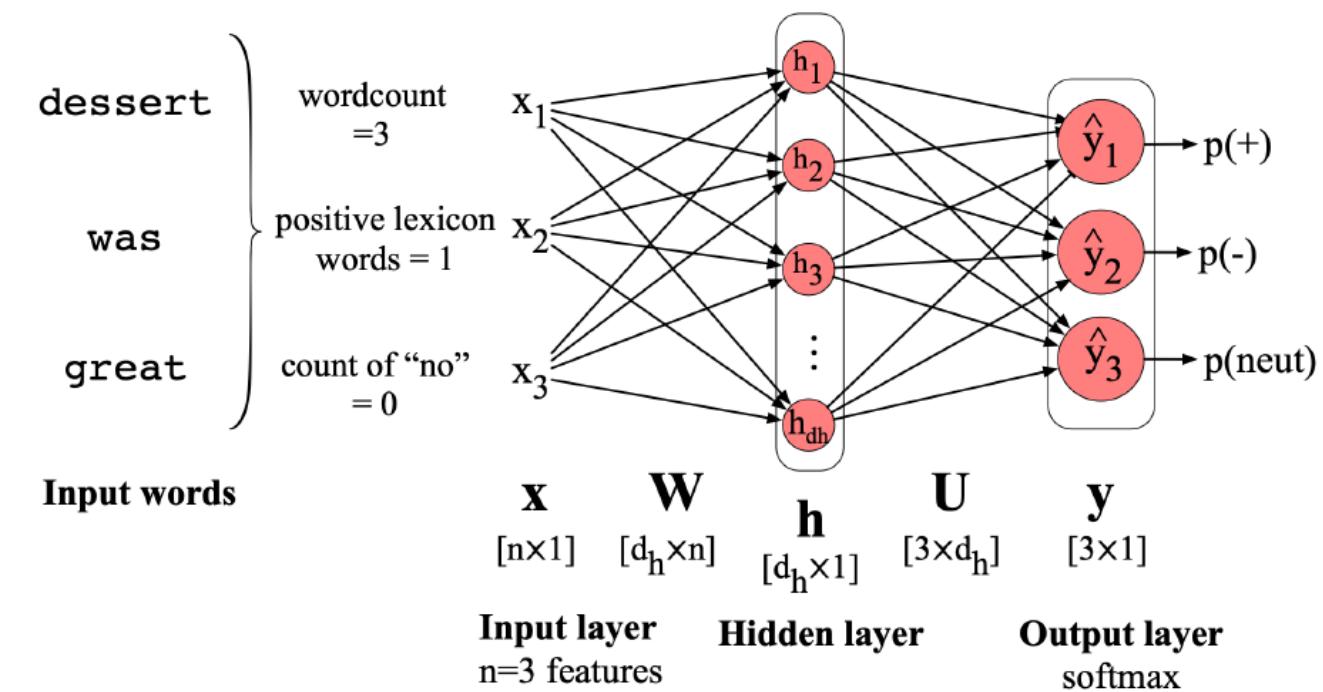
a gorgeous , witty , seductive movie .

label = positive

- Images: fixed-size input, continuous values
- Text: **variable-length** input, discrete words

# Solution 1: feature vector as input

- Input:  $w_1, w_2, \dots, w_K \in V$
- Output:  $y \in C$
- Input: dessert was great
- Output: positive  $C = \{\text{positive}, \text{negative}, \text{neutral}\}$



(each  $\mathbf{x}_i$  is a hand-designed feature)

- $\mathbf{x} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n]$
- $\mathbf{h} = \text{ReLU}(\mathbf{Wx} + \mathbf{b})$
- $\mathbf{y} = \mathbf{Uh}$
- $\hat{\mathbf{y}} = \text{softmax}(\mathbf{y})$

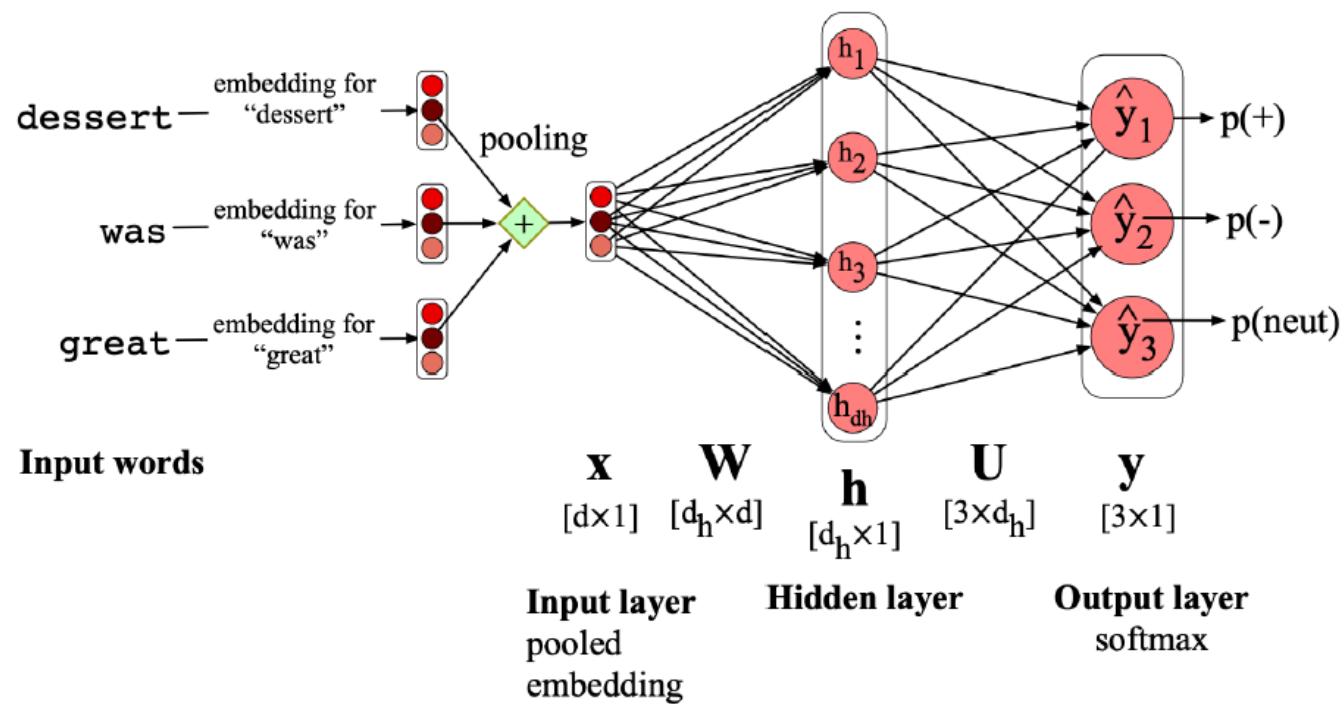
Deep learning has the promise to learn good features automatically..

# Solution 2: pooling on word embeddings

- Aggregate all the word embeddings into a vector via pooling function

$$\mathbf{x}_{\text{mean}} = \frac{1}{K} \sum_{i=1}^K e(w_i)$$

**pooling:** sum, mean or max



- $\mathbf{x} = \frac{1}{K} \sum_{i=1}^K e(w_i)$
- $\mathbf{h} = \text{ReLU}(\mathbf{W}\mathbf{x} + \mathbf{b})$
- $\mathbf{y} = \mathbf{U}\mathbf{h}$
- $\hat{\mathbf{y}} = \text{softmax}(\mathbf{y})$

Important note: each input has a different K

# Neural “bag-of-words” Model

---

Pro:

- This provides a simple and flexible way to handle variable-length input
- Learns feature representation (word embedding) automatically from data
- It can generalize to similar inputs through word embeddings

Cons:

- The model throw away any sequential information of the text

# Neural “bag-of-words” Model – Training

---

- $\mathbf{x} = \frac{1}{K} \sum_{i=1}^K e(w_i)$
- $\mathbf{h} = \text{ReLU}(\mathbf{Wx} + \mathbf{b})$
- $\mathbf{y} = \mathbf{Uh}$
- $\hat{\mathbf{y}} = \text{softmax}(\mathbf{y})$
- Training data:  $\{(d^{(1)}, y^{(1)}), \dots, (d^{(m)}, y^{(m)})\}$
- Parameters:  $\{\mathbf{W}, \mathbf{b}, \mathbf{U}\}$
- Optimize these parameters using stochastic gradient descent!
- Word embeddings can be treated as parameters too!

$$\mathbf{E} \in \mathbb{R}^{|V| \times d}$$

# Neural “bag-of-words” Model – Training

---

- Common practice: initialize word embedding  $E$  using static pre-trained word embeddings (e.g. word2vec), and optimize them using SGD
- When the training data is small, don't treat  $E$  as parameters
- When the training data is very large, initialization does not matter much, can use random initialization for word embeddings

# Feedforward Neural Language Models

# N-gram vs neural language models

---

Language models: Given  $x_1, x_2, \dots, x_n \in V$ , the goal is to model:

$$P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_1, \dots, x_{i-1})$$

Bigram:  $P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_{i-1})$

Maximum likelihood estimate:

Trigram:  $P(x_1, x_2, \dots, x_n) = \prod_{i=1}^n P(x_i | x_{i-2}, x_{i-1})$        $P(\text{sat}|\text{the cat}) = \frac{\text{count}(\text{the cat sat})}{\text{count}(\text{the cat})}$

Limitations: cannot handle long histories

# N-gram vs neural language models

---

- The larger n, the number of possible n-grams grow exponentially
- The larger n, the frequency counts for a specific n-gram is more sparse

$$P(w \mid \text{students opened their}) = \frac{\text{count(students opened their } w\text{)}}{\text{count(students opened their)}}$$

- A lot of contexts are similar and simply counting their frequency cannot reflect their semantic similarity

I am a **good** \_\_\_\_\_ count(I am a good w)

I am a **great** \_\_\_\_\_ count(I am a great w)

**e(good) ≈ e(great)**

Neural LM mitigates these issues

# Feedforward Neural Language Models

---

- Key Idea: instead of estimating raw probabilities, let's use a **feedforward neural network** to fit **the probabilistic distribution of language**

$$P(w \mid \text{I am a good})$$

$$P(w \mid \text{I am a great})$$

- Feedforward neural language models approximate the probability based on the previous  $m$  (e.g. 5) words –  $m$  is a hyper-parameter

$$P(x_1, x_2, \dots, x_n) \approx \prod_{i=1}^n P(x_i \mid x_{i-m+1}, \dots, x_{i-1})$$

# Feedforward Neural Language Models

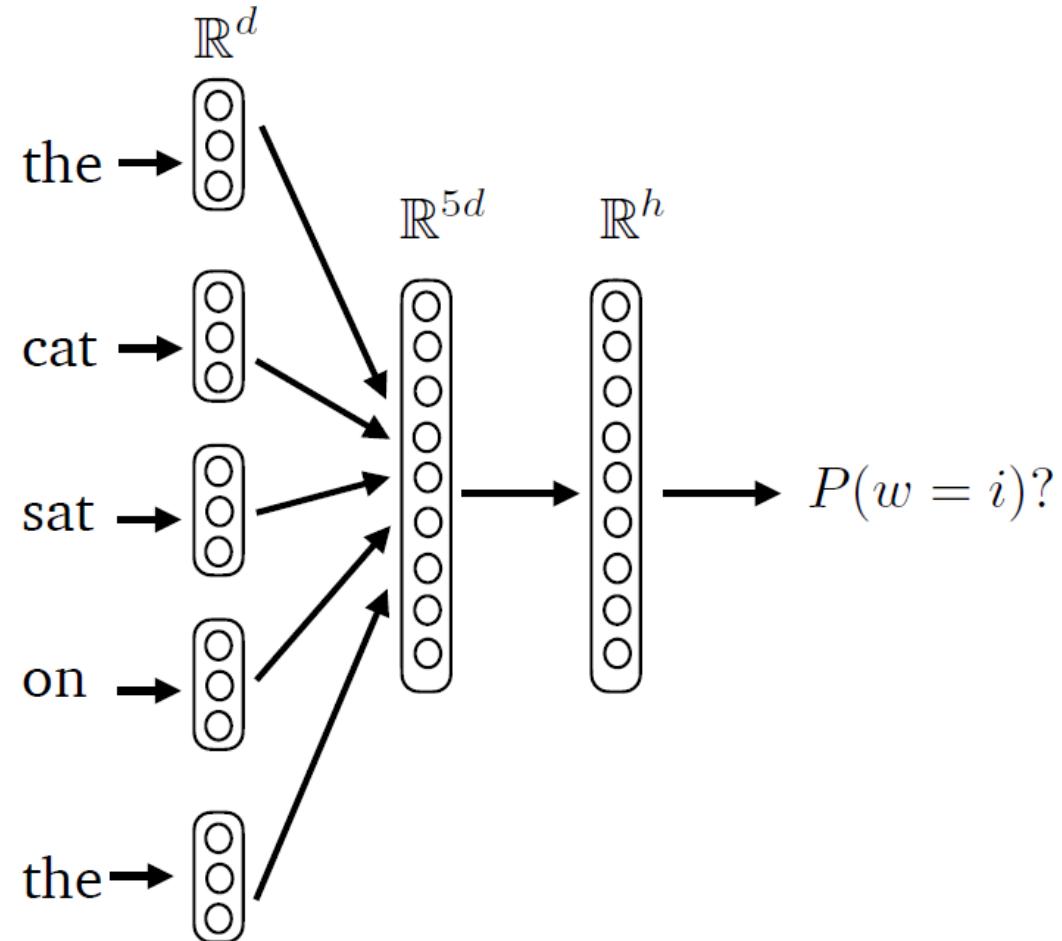
$$P(x_1, x_2, \dots, x_n) \approx \prod_{i=1}^n P(x_i | x_{i-m+1}, \dots, x_{i-1})$$

$P(\text{mat} | \text{the cat sat on the}) = ?$

d: word embedding size

h: hidden size

It is a  $|V|$ -way classification problem!



# Feedforward Neural Language Models

$P(\text{mat} \mid \text{the cat sat on the}) = ?$     d: word embedding size    h: hidden size

- Input layer (m= 5):    Q: why concat instead of taking the average?

$$\mathbf{x} = [e(\text{the}); e(\text{cat}); e(\text{sat}); e(\text{on}); e(\text{the})] \in \mathbb{R}^{md}$$

- Hidden layer:

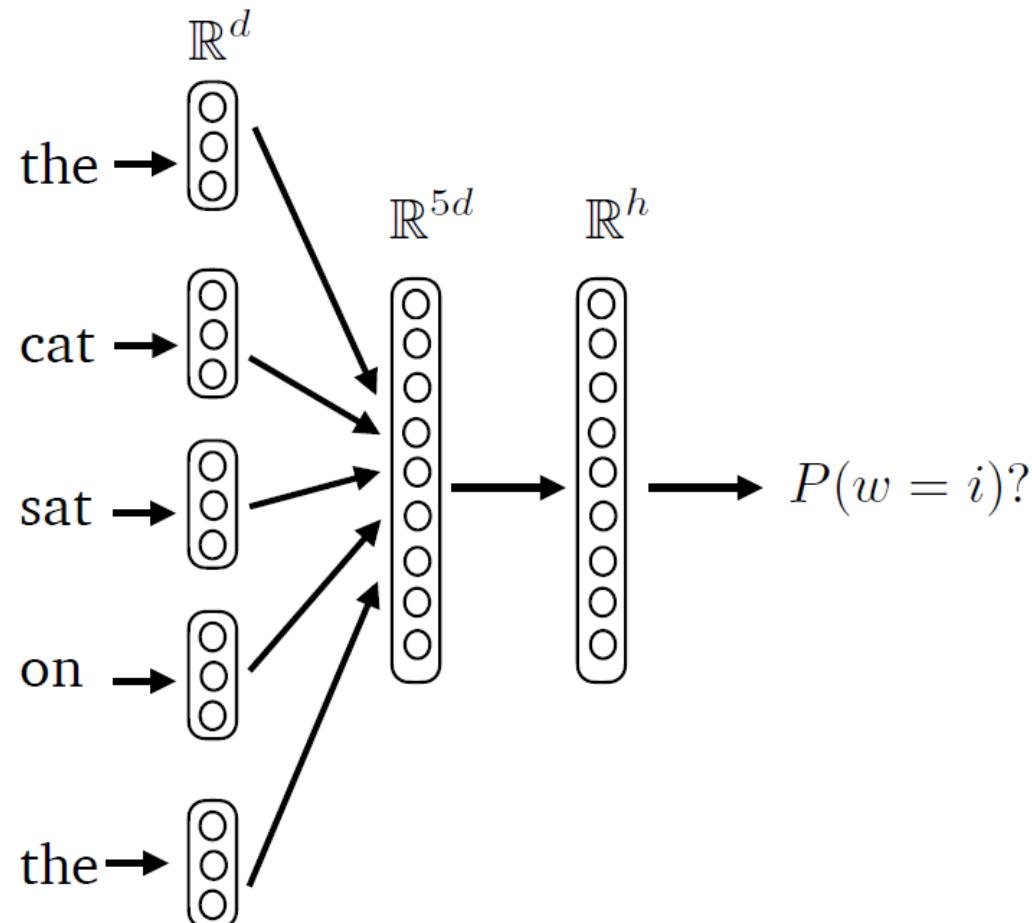
$$\mathbf{h} = \tanh(\mathbf{Wx} + \mathbf{b}) \in \mathbb{R}^h$$

- Output layer

$$\mathbf{z} = \mathbf{Uh} \in \mathbb{R}^{|V|}$$

$$P(w = i \mid \text{the cat sat on the})$$

$$= \text{softmax}_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_k e^{z_k}}$$



# Feedforward Neural Language Models

d: word embedding size, h: hidden size

- (a)  $\mathbf{W} \in \mathbb{R}^{h \times d}, \mathbf{U} \in \mathbb{R}^{|V| \times h}$
- (b)  $\mathbf{W} \in \mathbb{R}^{h \times 5d}, \mathbf{U} \in \mathbb{R}^{|V| \times h}$
- (c)  $\mathbf{W} \in \mathbb{R}^{h \times 5d}, \mathbf{U} \in \mathbb{R}^{|V| \times d}$
- (d)  $\mathbf{W} \in \mathbb{R}^{h \times d}, \mathbf{U} \in \mathbb{R}^{d \times h}$

Correct: (b)

- Input layer (m= 5):

$$\mathbf{x} = [e(\text{the}); e(\text{cat}); e(\text{sat}); e(\text{on}); e(\text{the})] \in \mathbb{R}^{md}$$

- Hidden layer:

$$\mathbf{h} = \tanh(\mathbf{W}\mathbf{x} + \mathbf{b}) \in \mathbb{R}^h$$

- Output layer

$$\mathbf{z} = \mathbf{U}\mathbf{h} \in \mathbb{R}^{|V|}$$

$$P(w = i \mid \text{the cat sat on the})$$

$$= \text{softmax}_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_k e^{z_k}}$$

# Feedforward Neural Language Models

- How to train this model? Use a lot of raw text to create training examples and run gradient-descent optimization

The Fat Cat Sat on the Mat is a 1996 children's book by Nurit Karlin. Published by Harper Collins as part of the reading readiness program, the book stresses the ability to read words of specific structure, such as -at.



the fat cat sat on → the  
fat cat sat on the → mat  
cat sat on the mat → is  
sat on the mat is → a  
...

- Limitations: (1)  $W$  linearly scales with the context size  $m$  (2) The model learns separate patterns for different positions

“sat on” corresponds to different parameters in  $W$

the fat cat **sat on** → the  
fat cat **sat on** the → mat  
cat **sat on** the mat → is

- Better solutions: Recurrent NNs, Transformers

A large, modern building with a glass facade and a metal frame under construction or renovation.

Thank you!

**UF** | Herbert Wertheim  
College of Engineering  
UNIVERSITY *of* FLORIDA

---

LEADING THE CHARGE, CHARGING AHEAD